

Darknet-NNPACK on Raspberry Pi

Aug. 2019

Ando Ki, Ph.D.


adki@future-ds.com

Contents

- What is NNPACK
- Building NNPACK
- Build and Run Darknet with NNPACK
- Gprof profiling
- Run Darknet with OpenCV
- Run Darknet with OpenCV and USB-Cam

What is NNPACK (1/2)

- <https://github.com/Maratyszczka/NNPACK>
- <https://github.com/shizukachan/NNPACK>
- <https://github.com/digitalbrain79/NNPACK-darknet>



- NNPACK is an acceleration package for neural networks on multi-core CPUs
 - ▶ Built-in expert-tuned kernels with very high performance:
 - Fast Fourier transform; Winograd transform
 - Matrix-matrix multiplication (GEMM); Matrix-vector multiplication (GEMV)
 - Max-pooling.
 - ▶ Multi-threaded SIMD-aware implementations of neural network layers.
 - ▶ Implemented in C99 and Python without external dependencies.

3

What is NNPACK (2/2)

- NNPACK layers
 - ▶ Convolutional layer
 - Training-optimized forward propagation (nnp_convolution_output)
 - Training-optimized backward input gradient update (nnp_convolution_input_gradient)
 - Training-optimized backward kernel gradient update (nnp_convolution_kernel_gradient)
 - Inference-optimized forward propagation (nnp_convolution_inference)
 - ▶ Fully-connected layer
 - Training-optimized forward propagation (nnp_fully_connected_output)
 - Inference-optimized forward propagation (nnp_fully_connected_inference)
 - ▶ Max pooling layer
 - Forward propagation, both for training and inference, (nnp_max_pooling_output)
 - ▶ ReLU layer (with parametrized negative slope)
 - Forward propagation, both for training and inference, optionally in-place, (nnp_relu_output)
 - Backward input gradient update (nnp_relu_input_gradient)
 - ▶ Softmax layer
 - Forward propagation, both for training and inference, optionally in-place (nnp_softmax_output)

4

Building NNPACK: prerequisites

- PIP
 - ▶ a tool for installing Python packages from the Python Package Index.
- PEACH-Py
 - ▶ Portable Efficient Assembly Code-generator in Higher-level Python
 - ▶ an assembler embedded in Python
- Confu
 - ▶ Ninja-based configuration system
- Ninja
 - ▶ a small build system with a focus on speed; e.g., Make
- NNPACK
 - ▶ <https://github.com/Maratyszczka/NNPACK>
 - ▶ <https://github.com/shizukachan/NNPACK>
 - ▶ Acceleration package for neural networks on multi-core CPUs

5

Building NNPACK: prepare for build

- Step 1: install pip if not available
 - ▶ `$ sudo apt-get install python-pip`
- Step 2: Install PeachyPy and Confu
 - ▶ `$ sudo pip install --upgrade git+https://github.com/Maratyszczka/PeachPy`
 - ▶ `$ sudo pip install --upgrade git+https://github.com/Maratyszczka/confu`
- Step 3: Install Ninja
 - ▶ `$ cd ~/work`
 - ▶ `$ git clone https://github.com/ninja-build/ninja.git`
 - ▶ `$ cd ninja`
 - ▶ `$ git checkout release`
 - ▶ `$./configure.py --bootstrap`
 - ▶ `$ export NINJA_PATH=$PWD`
 - ▶ `$ cd`

6

Building NNPACK: NNPACK

■ Step 4: Build NNPACK for Darknet (i.e., modified NNPACK for Darknet)

- ▶ \$ cd ~/work
- ▶ \$ git clone https://github.com/digitalbrain79/NNPACK-darknet.git
- ▶ \$ cd NNPACK-darknet
- ▶ \$ confu setup
- ▶ \$ python ./configure.py --backend auto
- ▶ \$ \$NINJA_PATH/ninja
- ▶ \$ sudo cp -a lib/* /usr/lib/
- ▶ \$ sudo cp include/nnpack.h /usr/include/
- ▶ \$ sudo cp deps/pthreadpool/include/pthreadpool.h /usr/include/
- ▶ \$ cd

7

Build Darknet with NNPACK

■ Step 5: Build Darknet

- ▶ \$ cd work
- ▶ \$ git clone https://github.com/digitalbrain79/darknet-nnpack.git
- ▶ \$ cd darknet-nnpack
 - ➡ Have a look at 'Makefile'
- ▶ \$ make

	Makefile
GPU=0	
CUDNN=0	
OPENCV=0	
NNPACK=1	
ARM_NEON=1	
OPENMP=0	
DEBUG=0	
....	

■ Step 6: Download weights

- ▶ \$ mkdir weights && cd weights
- ▶ \$ wget <https://pjreddie.com/media/files/yolov3-tiny.weights>
- ▶ \$ cd ..

ARM NEON is an advanced SIMD architecture for ARM processors which implement the ARMv7-A or ARMv7-R profile.

8

Run Darknet with NNPACK

Step 7: Run Darknet

► \$./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg

```
pi@raspberrypi: ~/work/codes/darknet-nnpack
File Edit Tabs Help
OPs
 22 conv    255  1 x 1 / 1    26 x  26 x 256    1 second with NNPACK & NEON
OPs
 23 yolo
Loading weights from weights/yolov3-tiny.weights...Done!
data/dog.jpg: Predicted in 1.033277 seconds.
dog: 57%
car: 52%
truck: 56%
car: 62%
bicycle: 59%
[pi@raspberrypi] █
```

```
pi@raspberrypi: ~/work/codes/darknet-nnpack
File Edit Tabs Help
 23 yolo
Loading weights from weights/yolov3-tiny.weights...Done!
data/dog.jpg: Predicted in 30.701036 seconds.
Bus error
[pi@raspberrypi] █
```

9

Gprof profiling

Compile with '-pg' option

```
ifeq ($(GPROF), 1)
CFLAGS+= -pg
endif
```

Run as normal

Run gprof with gmon.out

► \$ gprof darknet gmon.out

```
pi@raspberrypi: ~/work/codes/darknet-nnpack
File Edit Tabs Help
Flat profile:
Each sample counts as 0.01 seconds.
%   cumulative   self           calls     ms/call     ms/call    name
time  seconds    seconds
19.01    0.69    0.33  8845488      0.00      0.00  nnp_sgemm_only_4x12__neon
 7.60    0.82    0.13                581      0.15      0.17  compute_input_packing
 5.26    0.91    0.09                14      5.00      5.00  resize_image_compute_h
 4.68    0.99    0.08                13      5.38     30.38  activate
 4.09    1.06    0.07                760     0.07     0.24  copy_cpu
 4.09    1.13    0.07                760     0.07     0.07  make_convolutional_layer
 4.09    1.20    0.07                44      1.14     1.14  thread_main
 2.92    1.25    0.05                44      1.14     1.14  load_image_stb
 2.92    1.30    0.05                44      1.14     1.14  stbi_parse_zlib
 2.92    1.35    0.05                44      1.14     1.14  fill_cpu
```

10

Run Darknet with OpenCV

■ Step 5: Build Darknet with OpenCV

- ▶ ...
- ▶ \$ cd darknet-nnpack
- ▶ Have a look at 'Makefile'
- ▶ \$ make

GPU=0
CUDNN=0
OPENCV=1
NNPACK=1
ARM_NEON=1
OPENMP=0
DEBUG=0
....

Makefile

■ Step 6: Run Darknet

- ▶ \$./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg

11

Run Darknet with OpenCV and USB-Cam

■ Step 5: Build Darknet with OpenCV

- ▶ ...
- ▶ \$ cd darknet-nnpack
- ▶ Have a look at 'Makefile'
- ▶ \$ make

GPU=0
CUDNN=0
OPENCV=1
NNPACK=1
ARM_NEON=1
OPENMP=0
DEBUG=0
....

Makefile

■ Step 6: Run Darknet

- ▶ \$./darknet detector demo cfg/coco.data cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights

12

(주)퓨처디자인시스템

34051 대전광역시 유성구 문지로 193, KAIST 문지캠퍼스, F723호
(042) 864-0211~0212 / contact@future-ds.com / www.future-ds.com

Future Design Systems, Inc.

Faculty Wing F723, KAIST Munji Campus, 193 Munji-ro, Yuseong-gu, Daejeon 34051, Korea
+82-042-864-0211~0212 / contact@future-ds.com / www.future-ds.com



FUTURE
Design Systems